

Machine Learning Applied to Finance: Interest Rate Curve in Mexico

Jose Velasco

25/11/2020

Machine Learning Applied to Finance: Interest Rate Curve in Mexico

For the final capstone project of the HarvardX - PH125.9x, this paper examines the interest rates market in Mexico. Specifically, it models the 10-year rate in the swap interest rates market. Since different models are used, the accuracy of each model will be assessed using Root Mean Squared Error (RMSE).

1 Introduction

Interest rate markets are important to the economy because they shed light on economic matters. For example, by examining the overnight rate and the monetary policy statements of the different central banks, an idea can be formed about the current economic outlook (inflation and growth) and its expected course.

The economic theory also tell us the importance of the shape of the yield curve (commonly comprised of the overnight rate to the 10-year rate). As the US Central Bank (Federal Reserve or Fed for short) noted <https://www.chicagofed.org/publications/chicago-fed-letter/2018/404> the spread between the overnight rate and the long term rates are often regarded as predictors of recessions.

This paper will attempt to generate models from different approaches in order to obtain a sufficiently good prediction (measured using the RMSE loss function) for the 10yr rate.

First, an exploratory analysis will be conducted. Second, a series of models will be derived, starting from simple models (linear approaches) to machine learning ones (PCA, KNN and Random Forest). Third, the results will be presented comparing the different RMSEs. Last, conclusions and final remarks will be presented.

2 Methods and Analysis

2.1 Data Exploration

First of all, the data is loaded from the provided file. The information is extracted from Bloomberg.

```
library("readxl")
library("zoo")
library("forecast")
library("tidyverse")
library("dplyr")
library("TTR")
library("GGally")
library(tidyverse)
```

```
library(caret)
library(data.table)
library(stringr)
library(xts)
library(TSstudio)
```

```
#loading the dataset
dat <- read_excel("ratesmx.xlsx", sheet = 1, col_names = FALSE)
dat <- dat[-c(1,2,4,5,6),] #we remove useless rows
colnames(dat) <- dat[1,] #assign names to columns
colnames(dat)[1] <- "date"
names <- colnames(dat) #store names on a separate vector
dat <- dat[-c(1),] #remove names
dat <- as.matrix(dat) #transform to matrix
num_col <- ncol(dat) #extract the number of columns
num_row <- nrow(dat) # extract the number of rows
dat <- mapply(dat, FUN=as.numeric, nrow = length(dat)) #transform to numeric
#using extracted number of rows to preserve dimension
dat <- matrix(dat, ncol=num_col, nrow=num_row) #create a matrix again
#using extracted number of columns and rows to preserve dimension
dat <- data.frame(dat) # as data frame
dat[,1] <- as.Date(dat[,1], origin = "1899-12-30") #transform column one to dates
colnames(dat) <- names #restore names
#now we can work with the data
```

A quick summary of the data reveals the following:

```
summary(dat)
```

```
##      date      irs3m      irs6m      irs9m
## Min.   :2010-01-01  Min.   :3.250  Min.   :3.246  Min.   :3.270
## 1st Qu.:2012-09-04  1st Qu.:4.173  1st Qu.:4.245  1st Qu.:4.290
## Median :2015-05-08  Median :4.855  Median :4.881  Median :4.910
## Mean   :2015-05-08  Mean   :5.415  Mean   :5.437  Mean   :5.465
## 3rd Qu.:2018-01-09  3rd Qu.:7.289  3rd Qu.:7.228  3rd Qu.:7.135
## Max.   :2020-09-11  Max.   :8.635  Max.   :8.725  Max.   :8.795
##      irs1yr      irs2yr      irs3yr      irs4yr
## Min.   :3.313  Min.   :3.740  Min.   :4.165  Min.   :4.335
## 1st Qu.:4.332  1st Qu.:4.460  1st Qu.:4.760  1st Qu.:5.090
## Median :4.954  Median :5.115  Median :5.350  Median :5.580
## Mean   :5.501  Mean   :5.645  Mean   :5.832  Mean   :6.025
## 3rd Qu.:7.117  3rd Qu.:6.986  3rd Qu.:6.845  3rd Qu.:6.895
## Max.   :8.865  Max.   :8.880  Max.   :8.900  Max.   :8.925
##      irs5yr      irs7yr      irs10yr      overnight
## Min.   :4.459  Min.   :4.810  Min.   :5.120  Min.   :3.000
## 1st Qu.:5.350  1st Qu.:5.835  1st Qu.:6.221  1st Qu.:3.750
## Median :5.800  Median :6.230  Median :6.670  Median :4.500
## Mean   :6.212  Mean   :6.550  Mean   :6.864  Mean   :5.083
## 3rd Qu.:7.135  3rd Qu.:7.360  3rd Qu.:7.550  3rd Qu.:7.000
## Max.   :8.975  Max.   :9.095  Max.   :9.285  Max.   :8.250
##      fed      us10yr      tie28      mxn
## Min.   :0.2500  Min.   :0.5069  Min.   :3.274  Min.   :11.50
## 1st Qu.:0.2500  1st Qu.:1.8680  1st Qu.:4.066  1st Qu.:12.95
```

##	Median :0.2500	Median :2.2940	Median :4.840	Median :15.35
##	Mean :0.7265	Mean :2.2974	Mean :5.404	Mean :16.03
##	3rd Qu.:1.2500	3rd Qu.:2.7220	3rd Qu.:7.298	3rd Qu.:18.94
##	Max. :2.5000	Max. :3.9859	Max. :8.600	Max. :25.36
##	g1yr	g2yr	g3yr	g5yr
##	Min. :2.985	Min. :3.217	Min. :3.860	Min. :3.954
##	1st Qu.:3.956	1st Qu.:4.175	1st Qu.:4.705	1st Qu.:5.125
##	Median :4.558	Median :4.828	Median :5.167	Median :5.615
##	Mean :5.189	Mean :5.370	Mean :5.683	Mean :6.013
##	3rd Qu.:6.763	3rd Qu.:6.753	3rd Qu.:6.774	3rd Qu.:6.881
##	Max. :8.652	Max. :8.690	Max. :8.908	Max. :9.020
##	g10yr	g20yr	g30yr	
##	Min. :4.426	Min. :4.910	Min. :5.325	
##	1st Qu.:5.963	1st Qu.:6.634	1st Qu.:6.821	
##	Median :6.349	Median :7.106	Median :7.338	
##	Mean :6.596	Mean :7.164	Mean :7.368	
##	3rd Qu.:7.307	3rd Qu.:7.707	3rd Qu.:7.863	
##	Max. :9.257	Max. :9.793	Max. :9.838	

The data has more than 2,500 observations spanning from January 2010 to September 2020 (more than 10 years of daily observations). Furthermore, there are no missing values.


Here is a description of the variables:

- Index: dates
- irs3m: 3 months interest rate swap
- irs6m: 6 months interest rate swap
- irs9m: 9 months interest rate swap
- irs1yr: 1 year interest rate swap
- irs2yr: 2 years interest rate swap
- irs3yr: 3 years interest rate swap
- irs4yr: 4 years interest rate swap
- irs5yr: 5 years interest rate swap
- irs6yr: 7 years interest rate swap
- irs10yr: 10 years interest rate swap
- overnight: official Mexico central bank interest overnight rate (base rate)
- fed: official US central bank interest overnight rate (base rate)
- us10yr: maturity constant 10 years government bond rate
- tiie28: interbank one month rate for Mexico (base rate for swaps)
- mxn: exchange rate between us dollar and mexican peso
- g1yr to g30yr: constant maturity government bond rates for Mexico

The database provides two 10yr rates for Mexico: the IRS (interest rate swaps) and the government 10yr rate. The first one will be chosen for the model because irs rates present some advantages over government bonds. First of all, they are traded with constant maturity. This means that the observed rates are the actual level for a specific date on that instrument. In contrast, government bonds have specific maturity dates, for example, the current 10yr bond (or on the run bond) has a maturity date of May 2029, which means that it cannot be used as the 10yr bond if the analysis starts in 2010, since it would be the 20yr bond. For this reason, bloomberg calculates a theoretical 10yr government bond rate, taking into account the “on the run bond”, which changes from time to time. This means that the g1yr to g30yr are theoretical levels for that instruments. Hence, the choice of the IRS over the government bonds.

Regarding the IRS curve, we have information for 10 different nodes, spanning from 3 months to 10 years.


```
boxplot((dat[,2:11])) #boxplot of the IRS curve
```



The following visual examination shows the relationship between term and spread. The longer the maturity of the asset, the higher the premium vs the overnight rate. Note that this relation tends to hold under normal conditions. Under stress conditions or tightening monetary policy cycles (i.e. policy rate cuts) as in 2018-2019, the relationship tends to inverse.

```
#graph of the historical values of the IRS curve
dat %>%
  ggplot(aes(x = date)) +
  geom_line(aes(y = overnight), colour="#000000") +
  geom_line(aes(y = irs3m), colour="#2F4F4F") +
  geom_line(aes(y = irs6m), colour="#708090") +
  geom_line(aes(y = irs9m), colour="#778899") +
  geom_line(aes(y = irs1yr), colour="#696969") +
  geom_line(aes(y = irs2yr), colour="#808080") +
  geom_line(aes(y = irs3yr), colour="#A9A9A9") +
  geom_line(aes(y = irs4yr), colour="#C0C0C0") +
  geom_line(aes(y = irs5yr), colour="#0000FF") +
  geom_line(aes(y = irs7yr), colour="#6495ED") +
  geom_line(aes(y = irs10yr), colour="#00BFFF") +
  xlab("years") + ylab("rate (%)") +
  ggtitle("Interest Rate Swaps - Mexico")
```

Interest Rate Swaps – Mexico



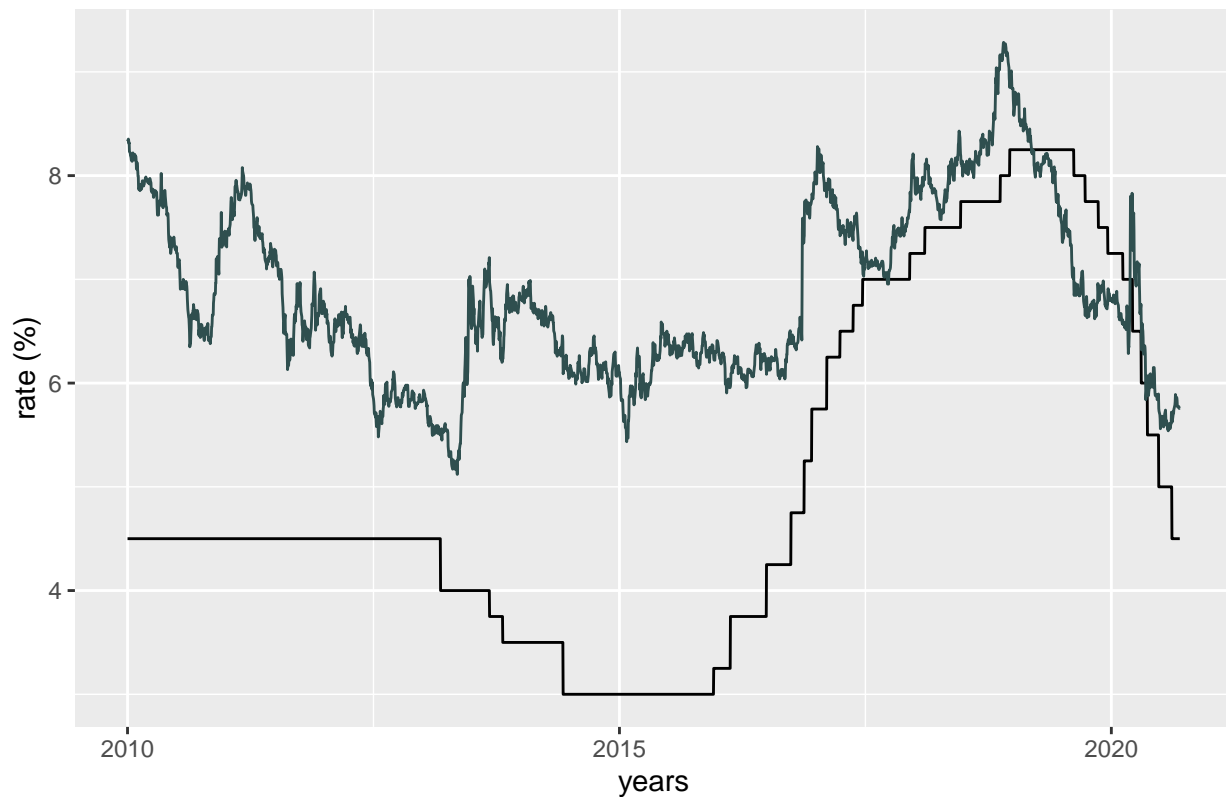
In the following graphs, a relationship can be observed between the 10yr IRS and three other variables.

First, it can be noted that the 10yr rate has a relation with the overnight rate. This makes sense as the one day rate is the base rate for the other ones as it is the one set by policy-makers.

```
#graph of the historical values of the overnight rate and the 10yr swap  
dat %>%
```

```
  ggplot(aes(x = date)) +  
  geom_line(aes(y = overnight), colour="#000000") +  
  geom_line(aes(y = irs10yr), colour="#2F4F4F") +  
  xlab("years") + ylab("rate (%)") +  
  ggtitle("Interest Rates (10yr and overnight) - Mexico")
```


Interest Rates (10yr and overnight) – Mexico



Second, a relationship appears between 10yrs from Mexico and US. Mexico's economy is dependent on US economy and so is the monetary policy. The Bank of Mexico watches closely US rates (and also inflation, growth, and other factors) in order to set the monetary policy rate.

```
dat %>%  
  ggplot(aes(x = date)) +  
  geom_line(aes(y = irs10yr, colour="#2F4F4F")) +  
  geom_line(aes(y = us10yr, colour="#778899")) +  
  xlab("years") + ylab("rate (%)") +  
  ggtitle("Interest Rates (10yr) - Mexico and US")
```


Interest Rates (10yr) – Mexico and US



Finally, a positive relation appears between the 10yr rate and the foreign exchange rate (USDMXN). Both variables are often regarded as a measure of risk for the Mexican assets. When the MXN appreciates against the dollar, it goes down in the graph. The same is true for the interest rate.

```
dat %>%
  ggplot(aes(x = date)) +
  geom_line(aes(y = irs10yr, colour="#2F4F4F")) +
  geom_line(aes(y = mxn, colour="#778899")) +
  xlab("years") +
  ggtitle("Interest Rates and FX - Mexico") +
  scale_y_continuous(
    name = "rate (%)",
    sec.axis = sec_axis(trans=~., name="mxn") #for the secondary axis
  )
```

Interest Rates and FX – Mexico



These relationships will be useful later when the models are derived.

2.2 Data Preparation

Now, the dataset is partitioned in training and dataset. The first one will comprise 80% of the dataset.

```
#we will create a data partition in training and testing sets  
#that is suitable for time series analysis, the train set comprises  
#the first 80% of the data time ordered  
  
set.seed(31416, sample.kind = "Rounding")  
test_index <- createDataPartition(dat$irs3m, times = 1, p = 0.2, list = FALSE) # create a 20% test set  
test_set <- dat[test_index,]  
train_set <- dat[-test_index,]
```

2.3 Model Evaluation

The model performance is going to be evaluated through the RMSE (residual mean square error). The function is defined as follows:

$$\text{RMSE} = \sqrt{\frac{1}{N} \sum_t (\hat{y}_t - y_t)^2}$$

Let y_t be defined as the rate at time t and denote our prediction with \hat{y}_t , with N being the number of time/node combinations and the sum occurring over all the series.

The RMSE is similar to a standard deviation: it is the typical error that is made when predicting a given node for a specific date. It translates in basis points (bps). For example, if the RMSE is 0.10, it means that the error made is from 10bps, which means an estimate of 5.50 instead of the true value of 5.60.

2.4 Base Model

After analysing the data, it is possible to define a base model. The one that will be the benchmark for the other, more complex, models.

Since we explored the relation between 10yr rates and overnight rate. A model which considers this plus a constant spread will be proposed. In this manner, the relationship could be explained with nothing more than a base rate, a constant.

Thus, this relationship can be modelled as follows:

$$\hat{y}_t = \beta_0 + \beta_{1,t}x_{1,t} + \varepsilon_t$$

Let \hat{y}_t be our estimate of the 10yr irs rate based on a constant β_0 , a linear term where $\beta_{1,t}$ is the effect derived from the overnight rate and ε_t is the error term due to randomness.

```
#base model
set.seed(1, sample.kind = "Rounding")
train_lm_1 <- train(irs10yr ~ overnight
                    , method = "lm", data = train_set)
#a linear model with only one variable and with intercept using the training set
```

The coefficients of the first model are presented:

```
#we observe the coefficients of the first model
train_lm_1[["finalModel"]][["coefficients"]]
```

```
## (Intercept)    overnight
##    5.1488317    0.3376162
```

The 10yr rate is basically explained by a fixed number and a spread vs the overnight rate.

Now, we extract the RMSE on the train set and observe a rather big error of 0.62. This means that the model does not do a good job predicting the 10yr interest rate. The error would be of 62 basis points (bps), which is huge for interest rates. However, now we have a starting point.

```
train_lm_1[["results"]][["RMSE"]] #displays the RMSE for the training set

## [1] 0.6214588
```

2.5 Economic Model

As we noted above, the 10-year Mexican rate is correlated to the US 10-year rates, to the overnight Mexican rate and to the USDMXN exchange rate. In this section we will derive a model that tries to capture this information:

We do not use the intercept since we are trying to explain the variable using exclusively other variables.

```
set.seed(1, sample.kind = "Rounding")
#linear model using 4 variables and without intercept
train_lm_2 <- train(irs10yr ~ overnight + fed + us10yr + mxn
, method = "lm", data = train_set,
tuneGrid = expand.grid(intercept = FALSE))
```

The coefficients of the second model are presented:

```
#displays the coefficients of the second model
train_lm_2[["finalModel"]][["coefficients"]]
```

```
## overnight      fed      us10yr      mxn
## 0.3560172 -0.5173115  1.2518984  0.1586797
```

In this model, the overnight Mexican rate, the 10yr US rate and exchange rate have a positive correlation with the 10yr IRS rate, whereas the Fed funds rate has a negative impact.

The RMSE from the train set is presented: 29bps, this model is twice as good as the base model according to the RMSE.

```
#get RMSE
train_lm_2[["results"]][["RMSE"]]
```

```
## [1] 0.2919127
```

2.6 Linear Model

In this model, the 10yr IRS rate is estimated using all the data available.

```
#liner model using all the variables and without intercept
set.seed(1, sample.kind = "Rounding")
train_lm_3 <- train(irs10yr ~ .
, method = "lm", data = train_set,
tuneGrid = expand.grid(intercept = FALSE))
```

The coefficients of the third model are presented:

```
#this will display the coefficients of all the variables
train_lm_3[["finalModel"]][["coefficients"]]
```

```
##      date      irs3m      irs6m      irs9m      irs1yr
## 2.660671e-07 -2.828782e-02  1.488440e-02  3.453592e-02 -1.418992e-02
##      irs2yr      irs3yr      irs4yr      irs5yr      irs7yr
## 8.534550e-02 -1.376729e-01 -1.026570e-02 -3.747744e-01  1.416160e+00
##      overnight      fed      us10yr      tiie28      mxn
## 4.361767e-03  5.094678e-02  2.462347e-02 -1.518824e-02  6.565777e-03
##      g1yr      g2yr      g3yr      g5yr      g10yr
## -1.085549e-02  2.239708e-02 -4.531583e-02 -4.305660e-02 -4.709839e-02
##      g20yr      g30yr
## 2.765568e-03  1.168624e-01
```

The train set RMSE has a substantial improvement to 2.45bps

```
#glance at the model
train_lm_3[["results"]][["RMSE"]]
```

```
## [1] 0.02453963
```

2.7 Principal Components Analysis

The fourth model that will be presented is a Principal Component Analysis (PCA) using only the IRS curve. The correlation between all the variables is presented here.

```
#correlation between all IRS
cor(train_set[,2:11])
```

```
##          irs3m    irs6m    irs9m    irs1yr    irs2yr    irs3yr    irs4yr
## irs3m    1.0000000 0.9973198 0.9920647 0.9837829 0.9535671 0.9182620 0.8825155
## irs6m    0.9973198 1.0000000 0.9985261 0.9939301 0.9706264 0.9392932 0.9057553
## irs9m    0.9920647 0.9985261 1.0000000 0.9983204 0.9808655 0.9531236 0.9218081
## irs1yr   0.9837829 0.9939301 0.9983204 1.0000000 0.9899089 0.9672499 0.9394269
## irs2yr   0.9535671 0.9706264 0.9808655 0.9899089 1.0000000 0.9930330 0.9768448
## irs3yr   0.9182620 0.9392932 0.9531236 0.9672499 0.9930330 1.0000000 0.9949320
## irs4yr   0.8825155 0.9057553 0.9218081 0.9394269 0.9768448 0.9949320 1.0000000
## irs5yr   0.8423381 0.8670265 0.8846882 0.9050067 0.9522917 0.9804330 0.9950531
## irs7yr   0.7740446 0.8000799 0.8194087 0.8427583 0.9016492 0.9425550 0.9701041
## irs10yr  0.7111108 0.7368423 0.7564341 0.7809931 0.8463054 0.8956391 0.9326913
##          irs5yr    irs7yr    irs10yr
## irs3m    0.8423381 0.7740446 0.7111108
## irs6m    0.8670265 0.8000799 0.7368423
## irs9m    0.8846882 0.8194087 0.7564341
## irs1yr   0.9050067 0.8427583 0.7809931
## irs2yr   0.9522917 0.9016492 0.8463054
## irs3yr   0.9804330 0.9425550 0.8956391
## irs4yr   0.9950531 0.9701041 0.9326913
## irs5yr   1.0000000 0.9888837 0.9622021
## irs7yr   0.9888837 1.0000000 0.9915305
## irs10yr  0.9622021 0.9915305 1.0000000
```

As we can see, the correlation between all the IRS is very high. We will limit the model to the IRS curve since an economic explanation can be applied to the first three components.

A dimension reduction algorithm is applied to the data.

```
#pca
x <- train_set[,2:11] %>% as.matrix() #we set the IRS curve as an x matrix
pca <- prcomp(x)
summary(pca)
```


```
## Importance of components:
##          PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation    4.1699 0.95156 0.25453 0.09338 0.02875 0.01776 0.01439
## Proportion of Variance 0.9466 0.04929 0.00353 0.00047 0.00005 0.00002 0.00001
## Cumulative Proportion 0.9466 0.99591 0.99944 0.99991 0.99996 0.99997 0.99998
##          PC8      PC9      PC10
```

```
## Standard deviation      0.01147 0.009645 0.008379
## Proportion of Variance 0.00001 0.000010 0.000000
## Cumulative Proportion  0.99999 1.000000 1.000000
```

It can be observed that the first three principal components explain 99.944% of the variation of the total data. As supported by the literature,

The PC1 one correspond to the level of the curve, which means that it approximates the funding rate (Banco de Mexico policy rate or overnight rate).

```
plot(pca$x[,1], train_set$overnight)
```




```
cor(pca$x[,1], train_set$overnight)
```

```
## [1] 0.9545597
```

The PC2 correspond to the slope of the curve. In this case, the slope between the 10yr and 2yr instruments is graphed along with the correlation. It is not as strong as the previous PC1 relation to the overnight rate.

```
plot(pca$x[,2], train_set$irs10yr - train_set$irs2yr)
```




```
cor(pca$x[,2], train_set$irs10yr - train_set$irs2yr)
```

```
## [1] 0.4601129
```

The PC3 correspond to the convexity of the curve. This one is traditionally approximated by a strategy called a fly, which implies buying the 2yr and 10yr instruments and selling the 5yr instrument. As the correlation and graph shows, it is not a good approximation of the PC3.

```
plot(pca$x[,3], train_set$irs2yr - train_set$irs5yr + train_set$irs10yr)
```



```
cor(pca$x[,3], train_set$irs2yr - train_set$irs5yr + train_set$irs10yr)
```

```
## [1] 0.02497903
```

2.7.1 Residuals In order to compare this model to the previous ones, a RMSE should be computed.

```
# we will extract the pcas and the rotation matrix
pc1to3 <- pca$x[,1:3]
pc_rot <- pca$rotation[,1:3]

#create a function to compute the reconstructed curve with the pca
curvapca <- function(i) {
  rowSums(pc1to3[i,]*(pc_rot))+pca$center
}

# create a matrix of the curve
curva_con_pca <- data.frame(t(sapply(seq(1:nrow(train_set)), curvapca)))


#then we can compute the squared residuals
pca_rmse <- sqrt(colSums((x - curva_con_pca)^2)/nrow(train_set))
pca_rmse
```

```
##      irs3m      irs6m      irs9m      irs1yr      irs2yr      irs3yr      irs4yr
```

```
## 0.04627066 1.15811065 2.61059092 0.04238561 3.11366575 0.64873808 0.03373044
##      irs5yr      irs7yr      irs10yr
## 1.02295273 1.01195821 0.04163830
```

```
#now we compute the dispersion of the data
dispersion <- (x - curva_con_pca)

#this show us the residuals
boxplot(dispersion)
```



As can be noted in the boxplot of the residuals and in the computation of the residual mean square error (RMSE), there is large variability in the following nodes: 6m, 9m, 2yr, 5yr and 7yr.

However, the model does provide a good estimate for the 10yr node. This could be explained by various reasons, but one of them is that the 10yr is a benchmark node and hence it is highly liquid in the market, whereas other nodes are not so liquid. The RMSE in the training set is about 4bps.

```
pca_rmse[10] #display only the residual por the 10yr irs
```

```
##      irs10yr
## 0.0416383
```


2.8 K Nearest Neighbours

The next model we will implement is a K-Nearest Neighbours approach.

```
# we will extract the pcas and the rotation matrix
set.seed(1, sample.kind = "Rounding")
train_knn <- train(irs10yr ~ ., method = "knn", data = train_set,
                  tuneGrid = data.frame(k = seq(1, 11, 1)))
train_knn$bestTune
```

```
##      k
## 2 2
```

```
plot(train_knn)
```



In order to get the best fit, cross-validation is used to find that when we set $k = 2$ we minimise our loss function to 7.5bps in the training set.

```
#we minimise the RMSE
min(train_knn[["results"]][["RMSE"]])
```

```
## [1] 0.07527954
```

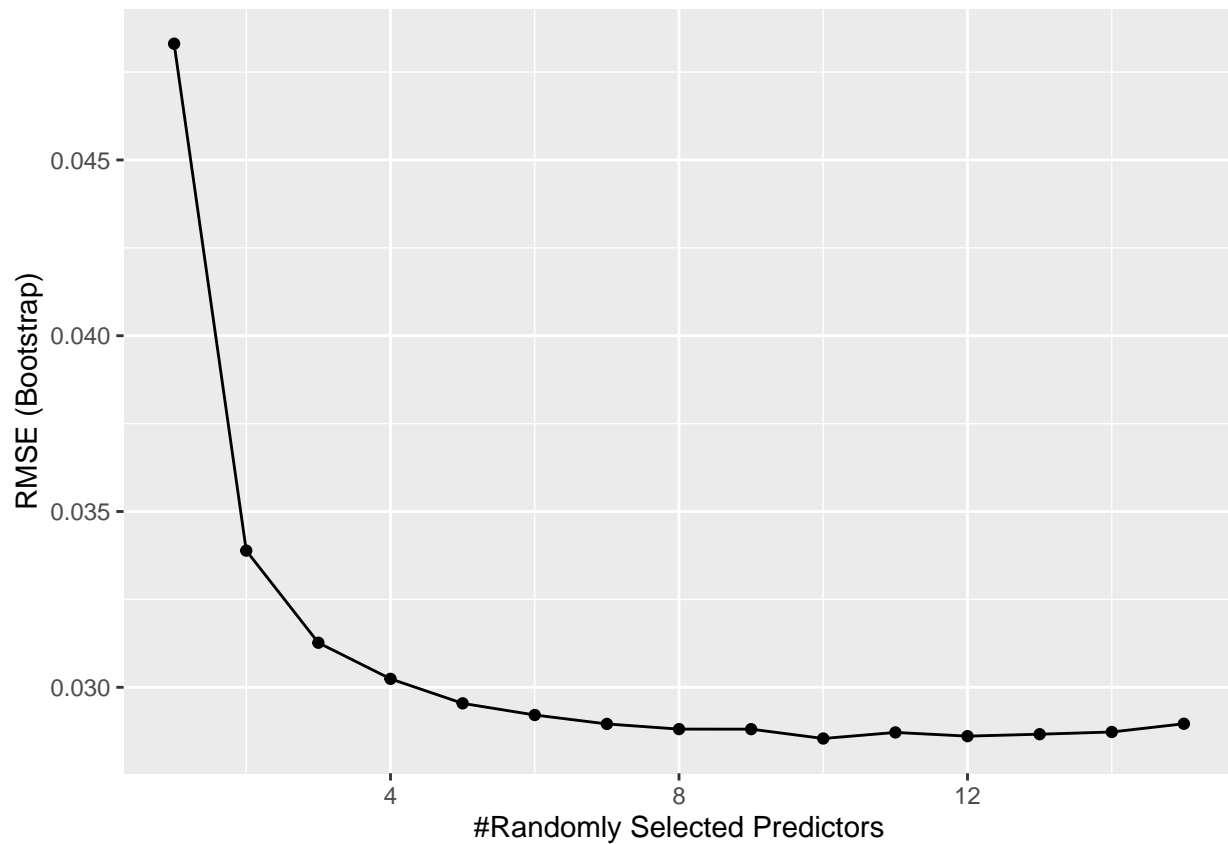
2.9 Random Forest

Finally, a random forest model is implemented.


```
####random forest#####
set.seed(1, sample.kind = "Rounding")
train_rf <- train(irs10yr ~ ., method = "rf", data = train_set,
                  tuneGrid = data.frame(mtry = seq(1, 15, 1)), ntree = 100)
train_rf$bestTune
```

```
##      mtry
## 10    10
```

```
ggplot(train_rf)
```



```
#and get the best rmse
min(train_rf[["results"]][["RMSE"]])
```

```
## [1] 0.02854399
```

As can be seen, a `mtry = 10` is chosen to minimise the loss function, which yields an RMSE in the training set of 2.85bps.

3 Validation and Results

After carefully estimate the models and calibrate them, it is possible to start the validation process using the test set. In this manner, the models can be compared in term of RMSE and a model can be chosen.

For the first model (baseline model), predictions can be derived from the following manner:

```
#correlation between all IRS
y_hat_lm_1 <- predict(train_lm_1, test_set, type = "raw")
```

Then, an RMSE is calculated for this baseline model using the test set.

```
#this would be for the validation
base_model_rmse <- sqrt(sum((test_set$irs10yr - y_hat_lm_1)^2)/nrow(test_set))
base_model_rmse
```


```
## [1] 0.6309902
```

The results are stored for comparison in the next table:

```
rmse_results <- tibble(method = "Baseline Model", RMSE = base_model_rmse)
```

It can be observed that the model does not do a good job fitting the data:

```
#this would be for the validation
plot(y_hat_lm_1, test_set$irs10yr)
```




For the second model (economic model) a similar process can be achieved. First, the predictions are estimated using the trained model and then an RMSE is estimated for the test set. A plot is also presented. Now, it can be noted that this model does a better job than the base model.

```
y_hat_lm_2 <- predict(train_lm_2, test_set, type = "raw")
second_model_rmse <- sqrt(sum((test_set$irs10yr - y_hat_lm_2)^2)/nrow(test_set))
second_model_rmse
```

```
## [1] 0.2844804
```

```
plot(y_hat_lm_2, test_set$irs10yr)
```



The results are stored:


```
rmse_results <- bind_rows(rmse_results,
  tibble(method="Economic Model", RMSE = second_model_rmse))
```

For the third model (linear model including all variables) a similar process can be achieved. First, the predictions are estimated using the trained model and then an RMSE is estimated for the test set. A plot is also presented.

```
y_hat_lm_3 <- predict(train_lm_3, test_set, type = "raw")
third_model_rmse <- sqrt(sum((test_set$irs10yr - y_hat_lm_3)^2)/nrow(test_set))
third_model_rmse
```

```
## [1] 0.02388543
```

```
plot(y_hat_lm_3, test_set$irs10yr)
```



The results are stored:

```
rmse_results <- bind_rows(rmse_results,
  tibble(method="Linear Model",
    RMSE = third_model_rmse))
```

For the fourth model (PCA model) a similar process can be achieved. First, the predictions are estimated using the trained model and then an RMSE is estimated for the test set. A plot is also presented.

```
#compute pca in test set
x <- test_set[,2:11] %>% as.matrix()
pca <- prcomp(x)

#computing residuals
# we will extract the pcas and the rotation matrix
pc1to3 <- pca$x[,1:3]
pc_rot <- pca$rotation[,1:3]


# create a matrix of the curve
curva_con_pca <- data.frame(t(sapply(seq(1:nrow(test_set)), curvapca)))

#then we can compute the squared residuals
pca_rmse <- sqrt(colSums((x - curva_con_pca)^2)/nrow(test_set))
```

```
fourth_model_rmse <- pca_rmse[10]
fourth_model_rmse
```

```
##      irs10yr
## 0.03962744
```

```
y_hat_pca <- as.numeric(unlist(curva_con_pca[10]))
plot(y_hat_pca, test_set$irs10yr)
```



The results are stored:


```
rmse_results <- bind_rows(rmse_results,
                          tibble(method="PCA Model",
                                RMSE = fourth_model_rmse))
```

For the fifth model (KNN) a similar process can be achieved. First, the predictions are estimated using the trained model and then an RMSE is estimated for the test set. A plot is also presented.

```
y_hat_knn <- predict(train_knn, test_set, type = "raw")
fifth_model_rmse <- sqrt(sum((test_set$irs10yr - y_hat_knn)^2)/nrow(test_set))
fifth_model_rmse
```

```
## [1] 0.05899655
```

```
plot(y_hat_knn, test_set$irs10yr)
```



The results are stored:


```
rmse_results <- bind_rows(rmse_results,  
  tibble(method="K Nearest Neighbours Model",  
    RMSE = fifth_model_rmse))
```

For the sixth model (RF) a similar process can be achieved. First, the predictions are estimated using the trained model and then an RMSE is estimated for the test set. A plot is also presented.

```
y_hat_rf <- predict(train_rf, test_set, type = "raw")  
sixth_model_rmse <- sqrt(sum((test_set$irs10yr - y_hat_rf)^2)/nrow(test_set))  
sixth_model_rmse
```

```
## [1] 0.02239061
```

```
plot(y_hat_rf, test_set$irs10yr)
```



The results are stored:

```
rmse_results <- bind_rows(rmse_results,
  tibble(method="Random Forest Model",
    RMSE = sixth_model_rmse))
```

This is the final summary of the results

```
## # A tibble: 6 x 2
##   method          RMSE
##   <chr>          <dbl>
## 1 Baseline Model 0.631
## 2 Economic Model 0.284
## 3 Linear Model   0.0239
## 4 PCA Model      0.0396
## 5 K Nearest Neighbours Model 0.0590
## 6 Random Forest Model 0.0224
```

method	RMSE
Baseline Model	0.6309902
Economic Model	0.2844804
Linear Model	0.0238854
PCA Model	0.0396274
K Nearest Neighbours Model	0.0589965
Random Forest Model	0.0223906

4 Conclusion

This paper has implemented six different models to estimate the 10 year IRS rate. Three of those models are linear models and three of them are machine learning models.

With respect to the linear models:

- The baseline model performs rather poorly (RMSE: 63bps) but it makes sense to estimate the 10 year rate with nothing but the overnight rate.
- The economic model uses theoretical relationships between interest rates, exchange rate, foreign interest rates and maturity. It yields better results but if implemented it would be used as a long-term model due to its short-term biases that tend to prevail during short term windows (RMSE: 28.4bps).
- The linear model that includes all the variables performs rather good (RMSE:2.4bps). Indeed it is the second best model. However, it requires a lot of variables as inputs. This is actually a problem for all the other models but the PCA. In terms of time of computing it is also efficient and in practice it would make an ideal model for short term estimates.
- The PCA model performs a little worst than the linear model (RMSE: 4bps) but it could be handy in practice as it only requires the IRS curve to provide a good estimate of the 10 year rate.
- The KNN model has an RMSE of 5.9bps. It does not outperform the linear model and it does require all the variables for its estimation. In consequence, it would not be the first choice when estimating rates.
- Finally, the Random Forest model outperforms all the other models with an RMSE of 2.2bps. However its time consuming and it is only slightly better than the linear model. Furthermore, it uses all the variables to achieve this RMSE. In practice it would be better to have a slightly worst model in terms of RMSE but quite faster, such as the PCA or the linear model.

Even if the Random Forest Model outperforms the other ones, in practice it would be useful to use more cost efficient models such as PCA or linear models. The models developed provide a statistical basis for estimation of the 10 year rate. It could further expand in the direction of a more theoretically sound model of the 10yr rate, including other variables such as inflation, equity index and so on.